
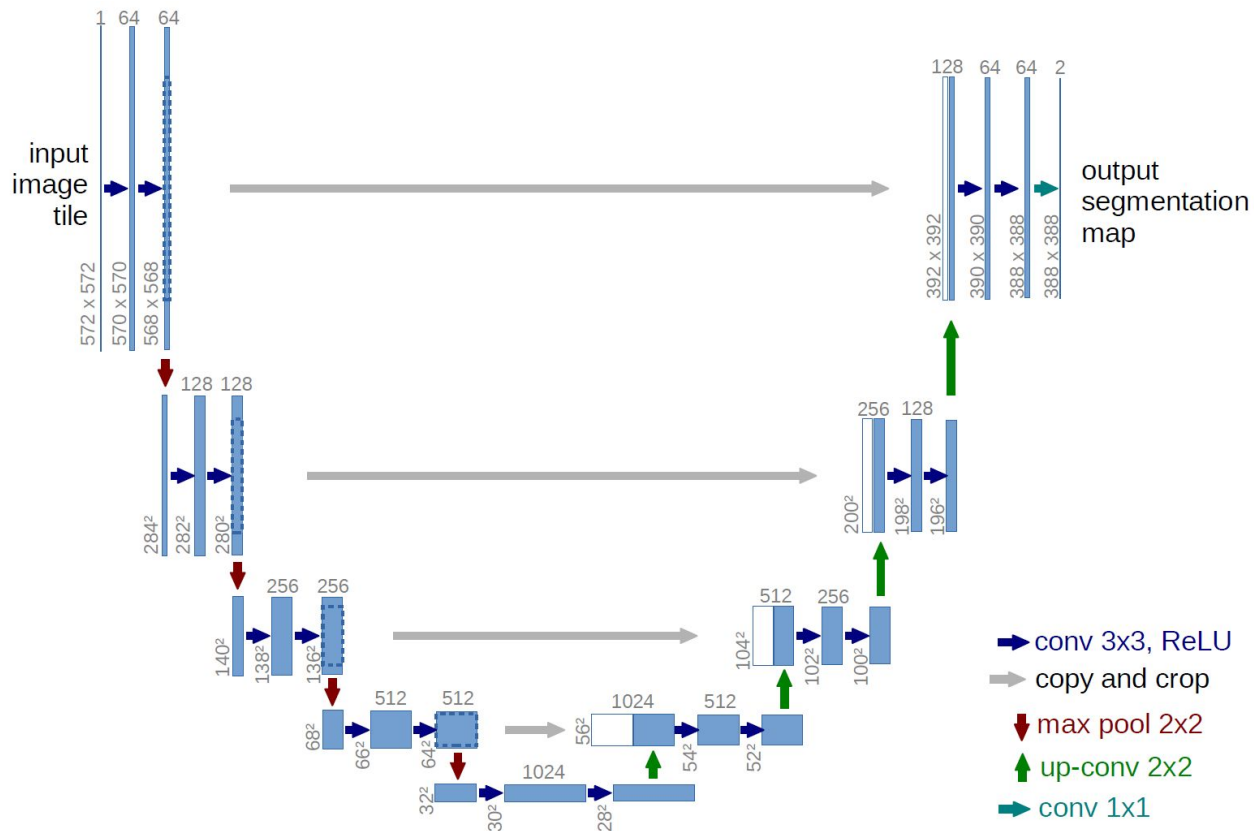


# ResNet50 based Image Compression

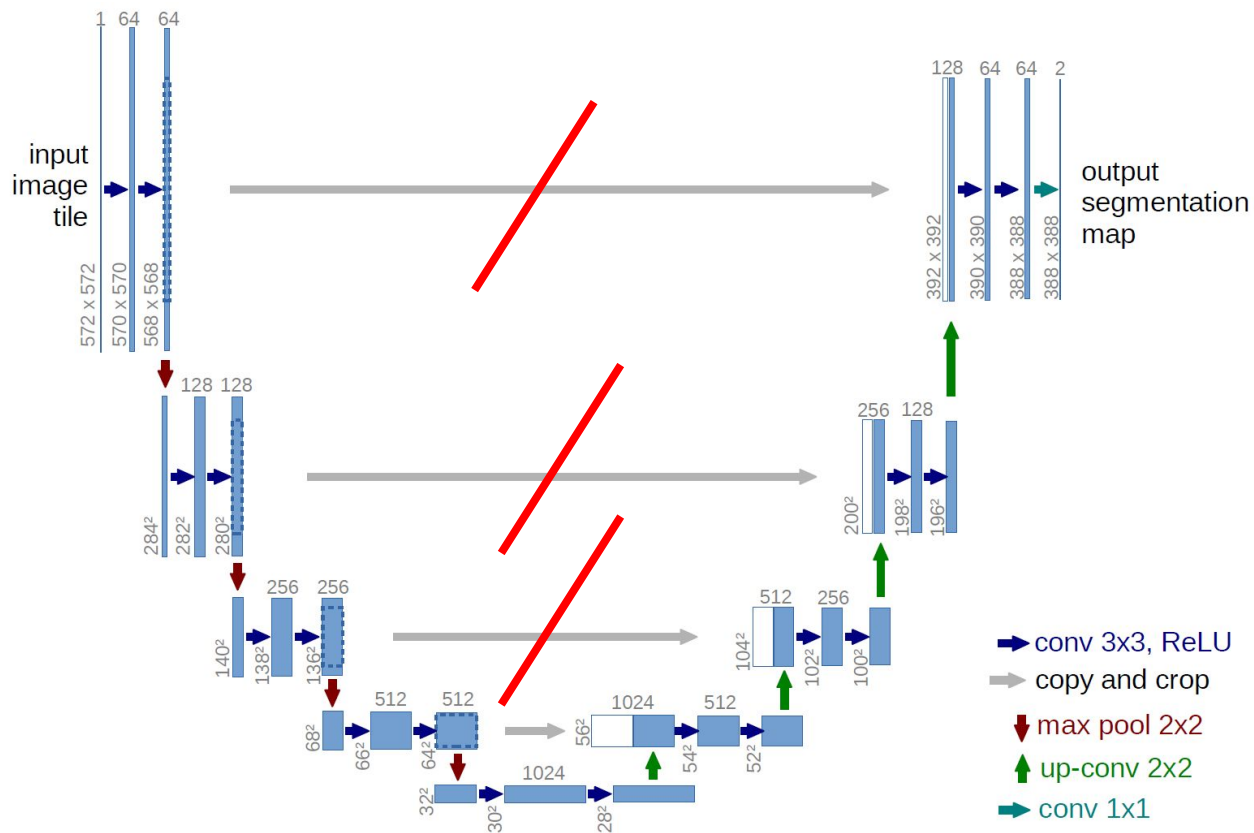


**нейро сжиматель**

# Изначальная идея (поток мыслей)



## Изначальная идея (не очень умно)



# Изначальная идея (не очень умно)

Old Forward

Pre-pool layer\_0 shape: torch.Size([2, 3, 512, 512]) total: **786432**

Pre-pool layer\_1 shape: torch.Size([2, 64, 256, 256]) total: **4194304**

Pre-pool layer\_2 shape: torch.Size([2, 256, 128, 128]) total: **4194304**

Pre-pool layer\_3 shape: torch.Size([2, 512, 64, 64]) total: **2097152**

Pre-pool layer\_4 shape: torch.Size([2, 1024, 32, 32]) total: **1048576**

Latent shape: torch.Size([2, 1024, 16, 16]) total: **262144**

Up block key: layer\_4 shape: torch.Size([2, 1024, 32, 32]) total: **1048576**

Up block key: layer\_3 shape: torch.Size([2, 512, 64, 64]) total: **2097152**


Up block key: layer\_2 shape: torch.Size([2, 256, 128, 128]) total: **4194304**

Up block key: layer\_1 shape: torch.Size([2, 64, 256, 256]) total: **4194304**

Up block key: layer\_0 shape: torch.Size([2, 3, 512, 512]) total: **786432**

Out shape: torch.Size([2, 3, 512, 512])

## Изначальная идея №2




Дальше Upscale или  
ConvTranspose2D

# Dataset

<https://www.kaggle.com/datasets/rhtsingh/130k-images-512x512-universal-image-embeddings>

**130k Images  
(512x512) -  
Universal Image  
Embeddings**



# Метрики

1. MSE
2. PSNR
3. Latent space Entropy (!)
4. Perceptual loss (VGG16)

*# MSE*

```
def mse_loss(result, target):  
    result = result.view(result.size(0), -1)  
    target = target.view(target.size(0), -1)  
    mses = F.mse_loss(result, target, reduction='none').mean(dim=1)  
    return torch.mean(mses)
```

*# PSNR*

```
def psnr(result, target):  
    result = result.view(result.size(0), -1)  
    maxes = torch.max(result, dim = 1)[0] ** 2  
    target = target.view(target.size(0), -1)  
    mses = F.mse_loss(result, target, reduction='none').mean(dim=1)  
    psnrs = 10.0 * torch.log10(maxes.to(torch.float32) /  
mses.to(torch.float32)).to(selected_dtype)  
    return torch.mean(psnrs)
```

```
def latent_entropy_aprox(result):  
    unique, counts = torch.unique(result, return_counts=True)  
    probabilities = counts.float() / result.numel()  
    entropy = -torch.sum(probabilities * torch.log2(probabilities + 1e-9)) #  
    return entropy.to(selected_dtype)  
  
# Perseptual Loss  
class VGGFeatures(nn.Module):  
    def __init__(self, feature_layers = {3, 6, 11, 16}):  
        super(VGGFeatures, self).__init__()  
        vgg = torchvision.models.vgg11(weights="VGG11_Weights.DEFAULT")  
        vgg.eval()  
        for param in vgg.parameters():  
            param.requires_grad = False  
        self.vgg = vgg  
        self.feature_layers = feature_layers  
    def forward(self, x):  
        features = []  
        for i, layer in enumerate(self.vgg.features):  
            x = layer(x)  
            if i in self.feature_layers:  
                features.append(x)  
        return features  
  
vgg_features = VGGFeatures().to(device)  
def vgg_perceptual_loss(result, target, vgg_to_use=vgg_features):  
    result = vgg_to_use(result)  
    target = vgg_to_use(target)  
    return sum(F.mse_loss(orig, decomp) for orig, decomp in zip(result, target))
```

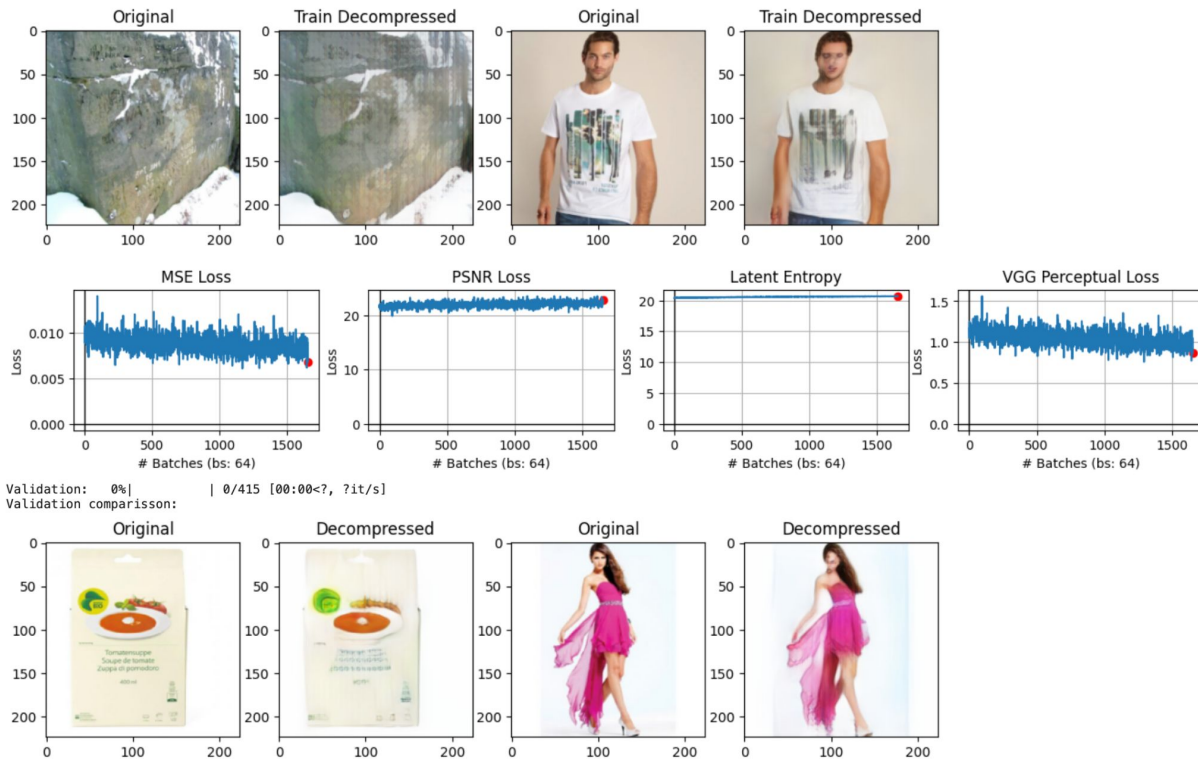
# Метрики

1. MSE (обучаем только на ней)

Считаем все, запоминаем


Потом расставляем коэф. на основе собранных данных

2. PSNR
3. Latent space Entropy (!)
4. Perceptual loss (VGG16)





# Казалось бы успех...



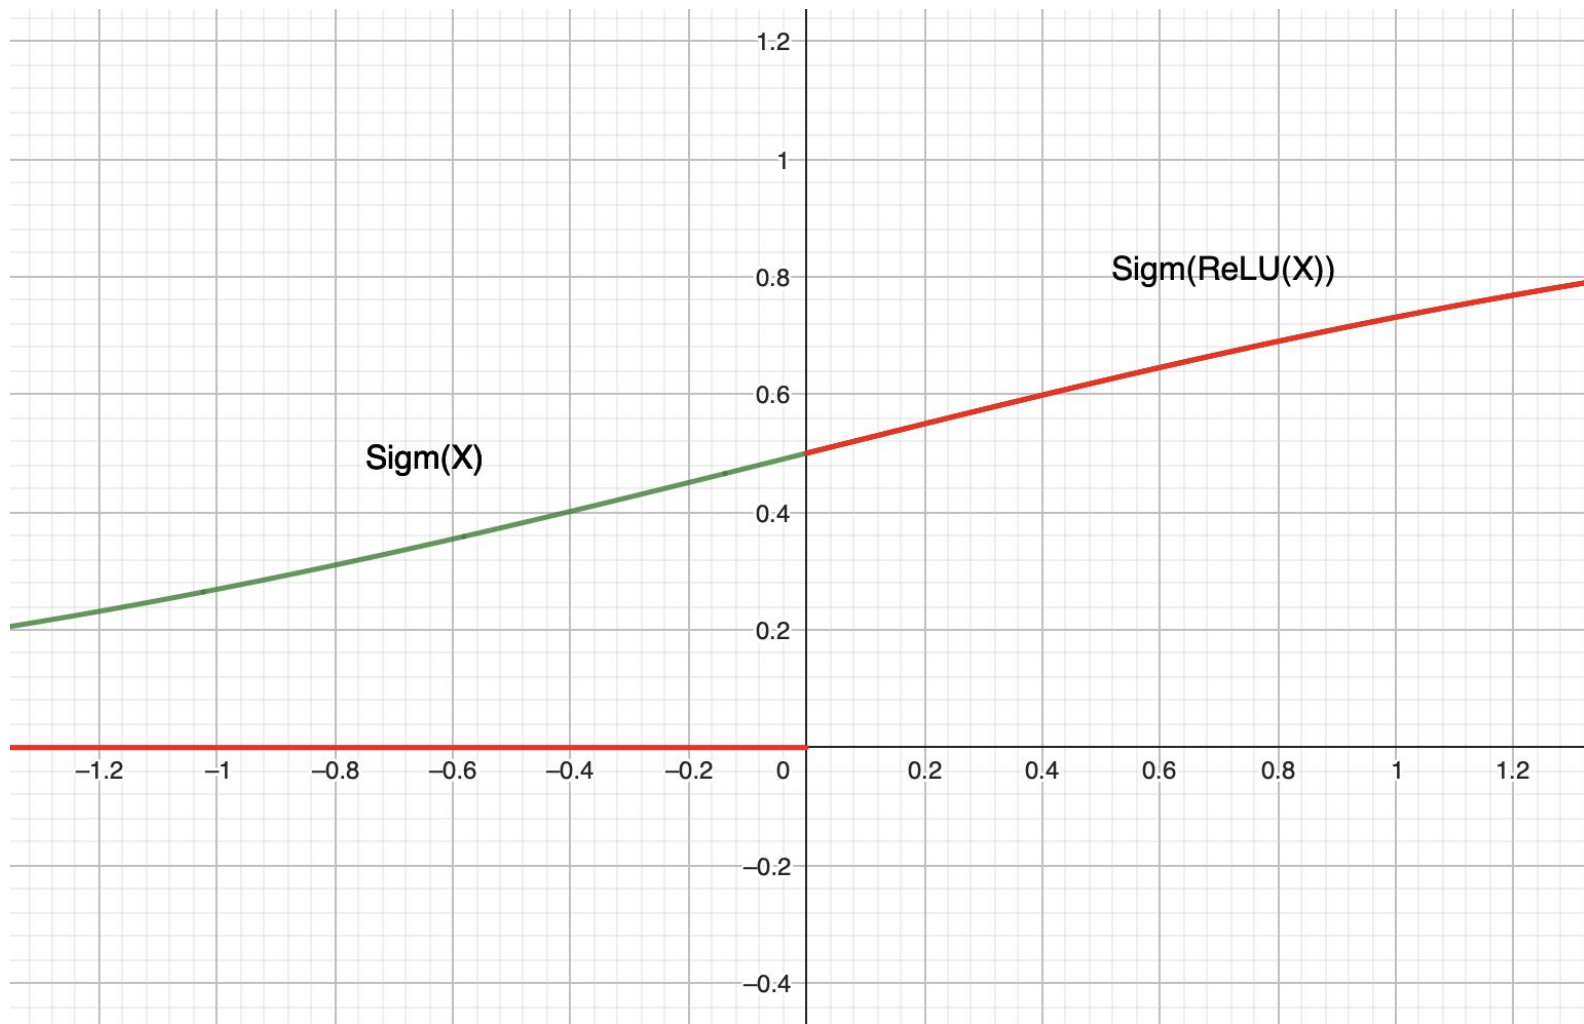
## Но что-то пошло не так

Промежуточное представление должно быть  $[0, 1]$

Функция активации: **sigmoid**

Однако перед этим шло **ReLU**...











# Результаты (Сравнение с JPEG)

